



DevOps at inBay Technologies

Part 2 • Continuous Delivery
John Marshall, Director Software Development




Goals


- Predictability
 - Velocity
 - Quality
- Agility
 - Short commitments, change lanes with minimal notice
 - Minimize the investment to validation cycle

Where we started

Development Cycle



Non-Stop Release Party



Aftermath




What we want to achieve



Develop



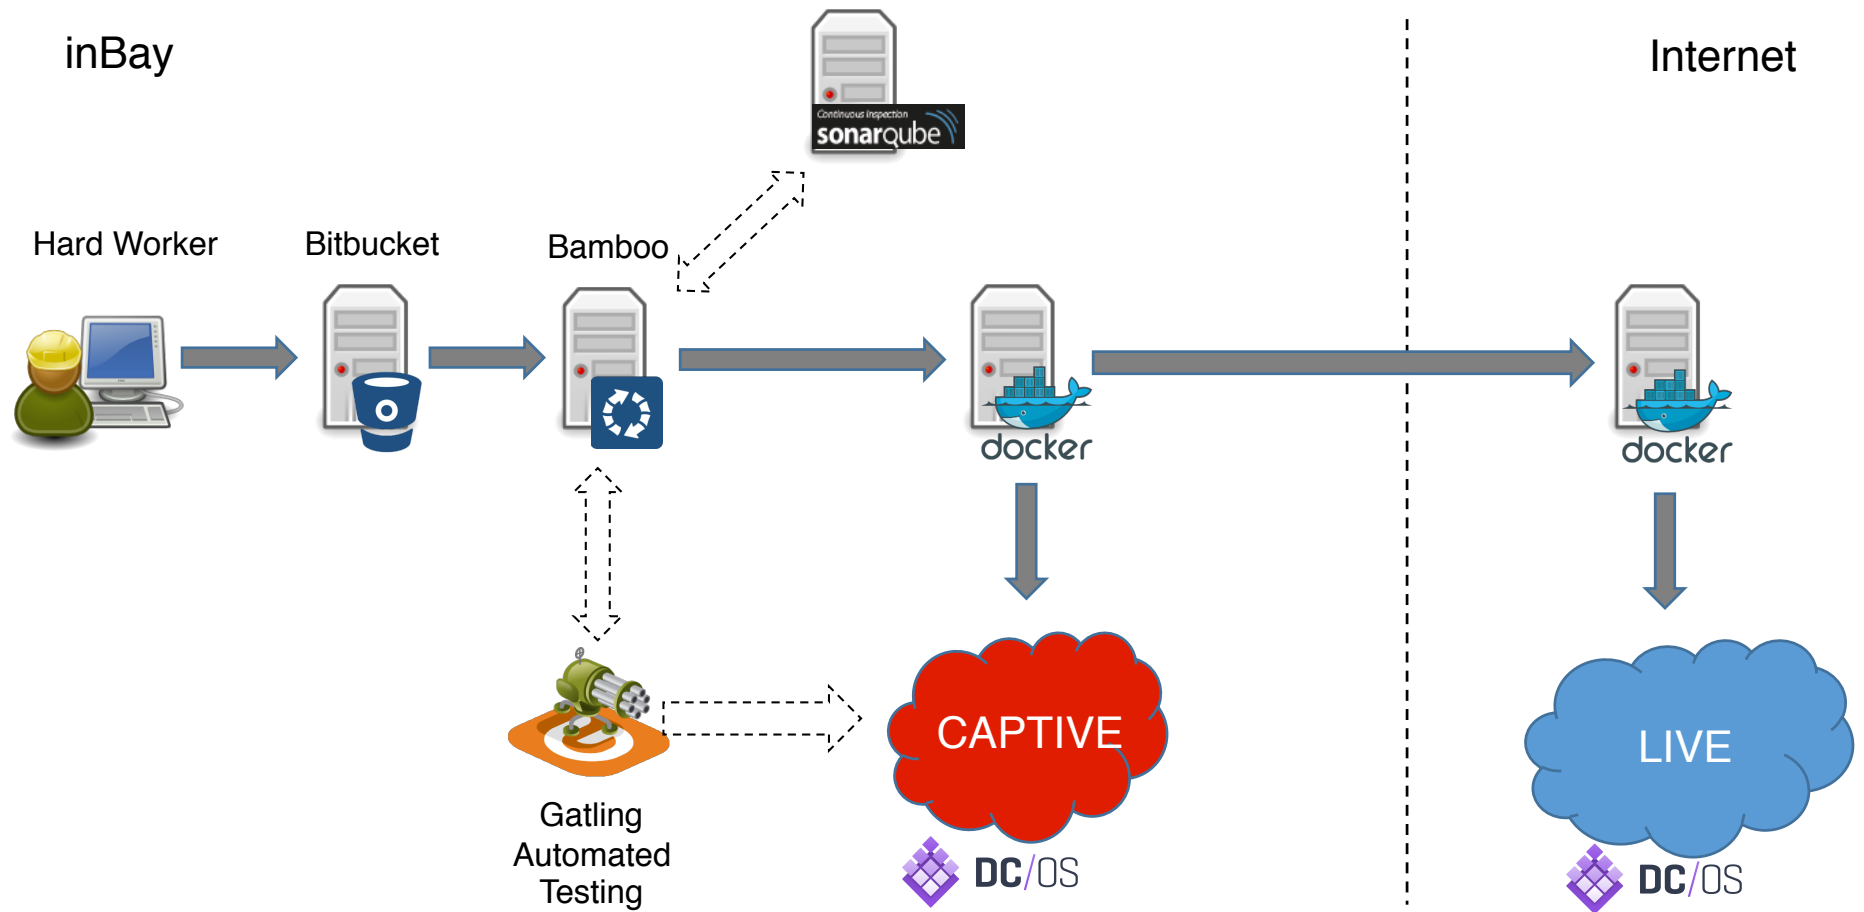
Build → Test → Deploy



Aftermath



Our Setup



Building...

Working with the CD Pipeline

- Requires a bit of a mind shift
 - Each submit will go live (on its own)
 - Treat the pipeline as a live system
 - Avoid application-test coupling
- Automate back pressure as well as forward flow
 - You need to have safeguards & road blocks
 - We use Sonarqube's quality gate
 - Breaks the build if a gate fails

Automated Versioning

- There are no snapshots
 - Every build is a potential release & must have a unique version
- Traditional versions like 1.1.0 don't really make sense
 - There are no 'releases' just a continual evolution of patches
- We use timestamps for versions
 - Dates are more meaningful than arbitrary version numbers
 - We include build metadata: Bamboo build number and git commit ID

House cleaning

- Docker artifacts accumulate in a hurry
 - On your development machine
 - On the build server
 - In the Docker Registry
 - In the cluster
- Automated cleanup is a necessity
 - We do this using scheduled jobs

Testing...

Predictability/agility only as good as your testing

- Start early (it's hard to catch up)
 - It will become part your velocity & culture
- Avoid accumulating test debt
 - You cannot reliably change uncovered code
 - You will end up writing these tests when you can least afford it
- Peer test reviews as/more important than code reviews
 - You can automate tests
 - But you need people to identify them
 - Best done before you implement

Invest in your test infrastructure


- You will write a lot of tests
 - Try and make the task as efficient and pleasant as possible
- We use Gatling, with it we have...
 - Build libraries of re-usable test steps, scenarios & support classes
 - Implemented environment variable based test configuration
 - Test suites can be run
 - Locally by developers via sbt
 - In the cloud in a Docker container

Mock External Services

- Prefer mocks over actual services, since mocks can
 - Perform message content validation
 - Inject error responses
- Our mock services
 - Have an automation API that the automated tests use
 - To validate operations
 - Inject behaviors

Deploying...

Our Environments




Live

- Production Services
- Behind HAProxy
- Secure Access
- Everything redundant



Testing

- Mock services
- Behind HAProxy
- Secure Access
- Everything redundant



Dev Lab

- Some real services
- Some services off
- Both secure and insecure access
- Non redundant



Local

- Services?
- No HAProxy
- Non-redundant
- Insecure access

Roll Your Own Dockers

- Allows you to have a common image hierarchy
- Common runtime & support environment
- Common place(s) to apply updates & patches
- More efficient Docker storage usage

Be environment independent

- All settings should be configurable by environment variables
 - Especially environment related settings
- Log to stdout
 - Use a Docker log driver to get the logs where you want them
- We can automate pushing the image through the pipe
Dev -> QA -> Staging -> Production
 - By supplying the relevant env var file at each stage
 - By allowing each environment to supply its own log routing

Configuration via Environment Variables

- We use Typesafe Config for our configuration
- We built a reference.conf preprocessor SBT plugin
 - Takes in src/.../reference.conf
 - Injects env var overrides for all the vars in the file
 - Outputs target/.../reference.conf
- We do this for important vendor files as well

Small company life

- You can't run a full cluster for every developer
 - Have a minimal deployment available for developer machines
 - Be able to run without things like https/tls
 - Private certificate authorities are a pain

Logging and troubleshooting with micro services

- Store logs in an indexed, searchable storage
 - We use Elasticsearch
 - Use a uniform log record layout
- Use correlators to link service spanning activities together
 - In inter-service messages
 - In log messages

Where did that Redis get to?

- In DC/OS things move around during deploys/restarts
 - Not everybody likes that, like things involving Redis
 - Your dev/test cluster should have space for things to move

Where we are today




Develop



???



Build → Test → Deploy



< 30 min



Aftermath

